

VERSION_1{BASE MODEL}:

Stock Prediction Model Using LSTM:

Introduction:

The goal of this Model is to build a predictive model for stock price trends of selected Indian banks. The model uses Long Short-Term Memory (LSTM) networks, a type of recurrent neural network suited to time-series data, to forecast whether the stock price will rise or fall.

Model Overview:

- **Objective:** Forecast stock price trends (increase or decrease).
- **Data Source:** Yahoo Finance.
- **Tools:** Python, yfinance, pandas, Keras, TA-Lib.
- **Approach:** LSTM neural networks with technical indicators as features.

Code Overview:

1. Setting Up the Environment

- **Importing Libraries:**
 - ``yfinance`` to fetch historical stock data.
 - ``pandas`` for data manipulation.
 - The ``bank_stocks`` list comprises the tickers of selected Indian banks.
- **Defining Time Range:** The period from 2013 to 2023, providing a decade of data for a comprehensive analysis.

2. Data Acquisition and Preprocessing

- **Fetching Stock Data:** Using ``yfinance.download()`` for each stock ticker.
- **Purpose:** To obtain historical stock data including open, close, high, low, and volume.

- **Handling Missing Data:** `fillna(method='ffill')` forward-fills missing values to maintain data continuity.
- **Saving Processed Data:** Exporting each stock's data frame to CSV for persistence and ease of access.

3. Feature Engineering

- **Technical Indicators:** Utilizing `ta.add_all_ta_features()` to add various technical indicators like RSI, MACD, and Bollinger Bands.
- **Rationale:** These indicators are crucial for analyzing market trends and momentum, which are predictive of future price movements.

4. Building the LSTM Model

- Creating a Sequential Model: To stack LSTM layers linearly.
- **Adding LSTM Layers:**
 - **First LSTM Layer:** `return_sequences=True` for passing sequential data to the next LSTM layer.
 - **Second LSTM Layer:** To further process the sequential information.
- **Output Layer:** A `Dense` layer with a `sigmoid` activation function, predicting the probability of a price increase.

5. Model Compilation and Training

- **Compilation:** Using `binary_crossentropy` for binary classification and `adam` optimizer for efficient learning.
- **Training:** Executing `model.fit()` to train the model over specified epochs and batch size.
- **Epochs:** The number of complete passes through the training dataset.
- **Batch Size:** The number of samples processed before the model is updated.

6. Model Evaluation

- **Predictions:** Using `model.predict()` to forecast stock trends on the test dataset.
- **Classification Report:** Assessing model performance with precision, recall, and f1-score.
- **Precision:** Proportion of true positive predictions.
- **Recall:** Ability of the model to find all positive samples.
- **F1-Score:** Harmonic mean of precision and recall.

Detailed Analysis of the Stock Prediction Model

1. Importing Libraries and Setting Parameters:

```
import yfinance as yf
import pandas as pd
```

- **yfinance:** Used for downloading historical stock data from Yahoo Finance.
- **pandas:** Essential for data manipulation and analysis.

```
bank_stocks = ['HDFCBANK.NS', 'SBIN.NS', 'ICICIBANK.NS',
               'AXISBANK.NS', 'KOTAKBANK.NS']
start_date = '2013-01-01'
end_date = '2023-12-25'
```

- **bank_stocks:** List of stock tickers for Indian banks. These represent the dataset's subjects.
- **start_date, end_date:** Define the time range for the historical data.

2. Downloading and Preprocessing Stock Data :

```
individual_stock_data = {}
for symbol in bank_stocks:
    print(f"Downloading data for {symbol}")
    stock_df = yf.download(symbol, start=start_date, end=end_date)
    stock_df.fillna(method='ffill', inplace=True)
    stock_df.to_csv(f'{symbol}.csv')
    individual_stock_data[symbol] = stock_df
```

- **For Loop:** Iterates through each stock ticker.
- **yf.download():** Fetches historical data for each ticker.
- **fillna():** Forward fills missing data points to maintain data integrity.
- **to_csv():** Saves each stock's processed data as a CSV file.
- **Dictionary Storage:** Each DataFrame is stored in a dictionary with the ticker as the key.

3. Technical Indicators with TA-Lib

```
import numpy as np
from ta import add_all_ta_features
from statsmodels.tsa.stattools import acf
```

- **TA-Lib:** Used for calculating various technical indicators.
- **acf:** Autocorrelation function from statsmodels, to analyze the correlation of the series with itself.

3. Adding Technical Indicators

```
df = add_all_ta_features(df, open="Open", high="High", low="Low",
close="Close", volume="Volume")
```

add_all_ta_features(): Adds multiple technical indicators to the DataFrame, like RSI, MACD, and Bollinger Bands, enhancing the dataset with metrics often used in stock market analysis.

4. LSTM Neural Network Model

```
from keras.models import Sequential
from keras.layers import LSTM, Dense
model = Sequential()
model.add(LSTM(units=50, return_sequences=True, input_shape=(1,
len(technical_indicators))))
model.add(LSTM(units=50))
model.add(Dense(1, activation='sigmoid'))
```

- **Sequential Model:** Basis for stacking layers in Keras.
- **LSTM Layer:** Captures time-dependency in data. **return_sequences=True** for the first LSTM layer to pass sequences to the next layer.
- **Dense Layer with Sigmoid Activation:** Outputs a probability prediction (between 0 and 1) indicating the likelihood of a stock price increase.

5. Model Compilation and Training

```
model.compile(loss='binary_crossentropy', optimizer='adam',
metrics=['accuracy'])
model.fit(X_train, y_train, epochs=100, batch_size=32, verbose=1)
```

- **compile()**: Configures the model with a loss function and optimizer.
- **binary_crossentropy** is suitable for binary classification.
- **fit()**: Trains the model.
- **epochs** define the number of times the model will work through the entire dataset.

6. Predicting and Evaluating the Model

```
yy_pred = model.predict(X_test)
classification_reports[symbol] = classification_report(y_test, y_pred,
output_dict=True)
```

- **predict()**: Generates output predictions.
- **classification_report**: Provides a detailed analysis of the model's accuracy, including precision, recall, and F1-score.

Classification Report for AXISBANK.NS:

	precision	recall	f1-score	support
0	0.449561	0.823293	0.581560	249.000000
1	0.443038	0.122378	0.191781	286.000000
accuracy	0.448598	0.448598	0.448598	0.448598
macro avg	0.446300	0.472835	0.386671	535.000000
weighted avg	0.446074	0.448598	0.373192	535.000000

Classification Report for HDFCBANK.NS:

	precision	recall	f1-score	support
0	0.481836	0.984375	0.646983	256.000000
1	0.666667	0.028674	0.054983	279.000000
accuracy	0.485981	0.485981	0.485981	0.485981
macro avg	0.574251	0.506524	0.350983	535.000000
weighted avg	0.578224	0.485981	0.338258	535.000000

Classification Report for ICICIBANK.NS:

	precision	recall	f1-score	support
0	0.440000	0.166667	0.241758	264.000000
1	0.494253	0.793358	0.609065	271.000000
accuracy	0.484112	0.484112	0.484112	0.484112
macro avg	0.467126	0.480012	0.425412	535.000000
weighted avg	0.467481	0.484112	0.427815	535.000000

Classification Report for KOTAKBANK.NS:

	precision	recall	f1-score	support
0	0.523220	0.628253	0.570946	269.000000
1	0.528302	0.421053	0.468619	266.000000
accuracy	0.525234	0.525234	0.525234	0.525234
macro avg	0.525761	0.524653	0.519783	535.000000
weighted avg	0.525747	0.525234	0.520069	535.000000

Classification Report for SBIN.NS:

	precision	recall	f1-score	support
0	0.000000	0.000000	0.000000	245.000000
1	0.542056	1.000000	0.703030	290.000000
accuracy	0.542056	0.542056	0.542056	0.542056
macro avg	0.271028	0.500000	0.351515	535.000000
weighted avg	0.293825	0.542056	0.381082	535.000000

The classification reports for each of the five bank stocks give us a detailed view of the performance of the LSTM models. Let's break down these results, focusing on the **precision**, **recall**, and **f1-score** for each class (**0** and **1**), as well as the **accuracy** and **macro avg** metrics.

1. AXISBANK.NS:

- The model has a high recall but low precision for class **0**, indicating it correctly identifies most of the actual **0** cases, but also incorrectly labels many **1**s as **0**s.
- For class **1**, both precision and recall are low, particularly recall, meaning the model struggles to correctly identify **1** cases.
- The overall accuracy is around 44.86%, and the macro average f1-score is about 38.67%, suggesting the model performs moderately but is biased towards predicting **0**.

2. HDFCBANK.NS:

- High recall but very low precision for class **0**, indicating a similar trend to AXISBANK.NS where the model predominantly predicts **0**.
- Extremely low recall for class **1**, showing the model almost never predicts **1** correctly.
- The accuracy is about 48.60%, but the model's effectiveness in distinguishing class **1** is poor, as reflected in the weighted average f1-score.

3. ICICIBANK.NS:

- Moderate recall for class **1** and low recall for class **0**. This model seems to be better at predicting **1** than **0**.
- The accuracy is close to 48.41%, with a macro average f1-score of about 42.54%, indicating a moderate performance with some bias.

4. KOTAKBANK.NS:

- More balanced precision and recall for both classes compared to other models.
- The accuracy is approximately 52.52%, with a macro average f1-score of around 51.98%. This suggests a relatively more balanced and better overall performance.

5. SBIN.NS:

- This model shows a peculiar behavior: it has high recall for class **1** but zero precision and recall for class **0**. It seems to predict **1** for almost all instances.
- The accuracy is 54.21%, but this is misleading as the model fails to identify any **0** cases correctly.

Conclusions:

- The models for AXISBANK.NS, HDFCBANK.NS, and ICICIBANK.NS show a bias towards predicting class **0** (with varying degrees), resulting in high recall but low precision for class **0**.
- KOTAKBANK.NS's model appears to be the most balanced among the five, with a more even distribution of precision and recall across classes.
- The model for SBIN.NS is heavily biased towards predicting class **1**, failing to identify any **0** cases.

- The accuracy and macro average f1-scores suggest moderate performance, but these models may need further tuning and possibly more complex architectures or additional features to improve their predictive power.
- The varying performances across different stocks highlight the unique characteristics and behaviors inherent in each stock's data.

Below are Improvements I am thinking of as of now:

- Experiment with different LSTM architectures or additional layers to capture more complex patterns.
- Consider using more or different technical indicators as features.
- Adjust the training parameters, like the number of epochs and batch size..
- Implement cross-validation to ensure the models' robustness.